

tsibble and fable



with xaringan

Kevin Song

2020-02-19

1) tsibble

2) fable

1) tsibble

```
install.packages("tsibble")
```

tsibble package provides a data infrastructure for tidy temporal data with wrangling tools adapting tidy data principles.

- tsibble is a data- and model-oriented object.
- In tsibble:
 1. **Key** is a set of variables that define observational units over time.
 2. **Index** is a variable with inherent ordering from past to present.
 3. Each observation should be uniquely identified by **key** and **index**.
 4. Each observational unit should be measured at a common **interval**, if regularly spaced.

Coerce to a tsibble - `as_tsibble()`

- To coerce data frame to tsibble, we need to declare **key** and **index**.
- The **key** can be comprised of empty, one, or more variables.
 - For multiple variables, declare `key = c(x1, x2, x3)`.
 - can be created in conjunction with tidy selectors like `starts_with()`
- The **Index** supports an extensive range of indices
 - native time classes in R (Date, POSIXct and difftime)
 - tsibble's new additions (yearweek, yearmonth, and yearquarter)
- The **Interval** is computed from index based on the representation.

Interval	Class
Annual	interger/double
Quarterly	yearquarter
Monthly	yearmonth
Weekly	yearweek
Daily (Subdaily)	Date / difftime (POSIXt / difftime / hms)

Ex) nycflights13::weather

- `time_hour` containing the date-times should be declared as **index**
- `origin` should be declared as **key**.
- Other columns can be considered as measured variables.

```
weather <- nycflights13::weather %>%  
  select(origin, time_hour, temp, humid, precip)
```

```
weather
```

```
## # A tibble: 26,115 x 5
```

##		origin	time_hour	temp	humid	precip
##		<chr>	<dtm>	<dbl>	<dbl>	<dbl>
##	1	EWR	2013-01-01 01:00:00	39.0	59.4	0
##	2	EWR	2013-01-01 02:00:00	39.0	61.6	0
##	3	EWR	2013-01-01 03:00:00	39.0	64.4	0
##	4	EWR	2013-01-01 04:00:00	39.9	62.2	0
##	5	EWR	2013-01-01 05:00:00	39.0	64.4	0
##	6	EWR	2013-01-01 06:00:00	37.9	67.2	0
##	7	EWR	2013-01-01 07:00:00	39.0	64.4	0
##	8	EWR	2013-01-01 08:00:00	39.9	62.2	0

Ex) nycflights13::weather

- `time_hour` containing the date-times should be declared as **index**
- `origin` should be declared as **key**.
- Other columns can be considered as measured variables.

```
weather_tsbl <- as_tsibble(weather,  
                             key = origin,  
                             index = time_hour)  
weather_tsbl
```

```
## # A tsibble: 26,115 x 5 [1h] <America/New_York>  
## # Key:          origin [3]  
##   origin time_hour      temp humid precip  
##   <chr>   <dtm>         <dbl> <dbl> <dbl>  
## 1 EWR    2013-01-01 01:00:00  39.0  59.4     0  
## 2 EWR    2013-01-01 02:00:00  39.0  61.6     0  
## 3 EWR    2013-01-01 03:00:00  39.0  64.4     0  
## 4 EWR    2013-01-01 04:00:00  39.9  62.2     0  
## 5 EWR    2013-01-01 05:00:00  39.0  64.4     0  
## 6 EWR    2013-01-01 06:00:00  37.9  67.2     0  
## 7 EWR    2013-01-01 07:00:00  39.0  64.4     0
```

fill_gaps()

| **fill_gaps()** turns implicit missing values into explicit.

- If the observations are made at regular time interval, we could turn these implicit missingness to be explicit simply using **fill_gaps()**.
- **fill_gaps()** also handles filling in time gaps by values or functions, and respects time zones for date-times.

```
full_weather <- weather_tsbl %>%  
  fill_gaps(precip = 0) %>%  
#filling gaps in precip with 0  
  group_by_key() %>%  
  tidyr::fill(temp, humid, .direction = "down")  
#replaces NAs with its previous observation for each origin (key)  
  
full_weather
```

```
## # A tsibble: 26,190 x 5 [1h] <America/New_York>  
## # Key:      origin [3]  
## # Groups:   origin [3]  
##   origin time_hour      temp humid precip  
##   <chr>   <dtm>         <dbl> <dbl> <dbl>
```

`index_by()` + `summarise()`

■ `index_by()` + `summarise()` to aggregate over calendar periods

- `index_by()` is the counterpart of `group_by()` in temporal context, but it groups the index only.
- In conjunction with `index_by()`, `summarise()` and its scoped variants aggregate variables over calendar periods.
- `index_by()` goes hand in hand with the index functions including `as.Date()`, `yearweek()`, `yearmonth()` and `yearquarter()`, as well as other friends from `lubridate`.

index_by() + summarise()

Example) To compute average temperature and total precipitation *per month*, we can apply `yearmonth()` to the index variable (referred to as `.`).

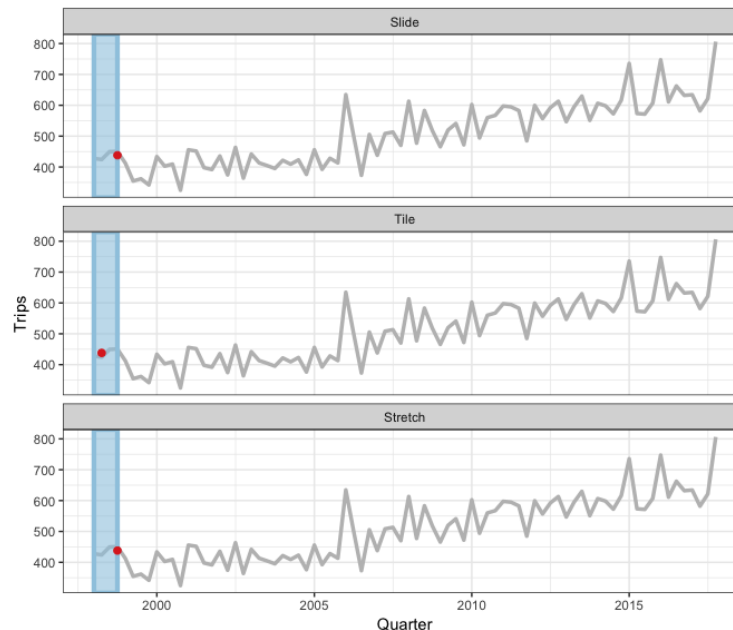
- While collapsing rows (by `summarise()`), `group_by()` and `index_by()` take care of updating the key and index respectively.

```
full_weather %>%  
  group_by_key() %>%  
  index_by(year_month = ~ yearmonth(.)) %>% # monthly aggregates  
  summarise(  
    avg_temp = mean(temp, na.rm = TRUE),  
    ttl_precip = sum(precip, na.rm = TRUE)  
  )
```

```
## # A tsibble: 36 x 4 [1M]  
## # Key:      origin [3]  
##   origin year_month avg_temp ttl_precip  
##   <chr>      <mth>      <dbl>      <dbl>  
## 1 EWR      2013 Jan      35.6        3.53  
## 2 EWR      2013 Feb      34.2        3.83  
## 3 EWR      2013 Mar      40.1         3  
## 4 EWR      2013 Apr      53.0        1.47
```

Rolling with `slide()`, `tile()`, `stretch()`

- Several functions in `tsibble` allow for different variations of moving windows using `purrr`-like syntax:
 - `slide()`: sliding window with overlapping observations.
 - `tile()`: tiling window without overlapping observations.
 - `stretch()`: fixing an initial window and expanding to include more observations.



2) fable

```
install.packages("fable")
```

fable package provides a collection of commonly used univariate and multivariate time series forecasting models

- Forecasting models include exponential smoothing via state space models and automatic ARIMA modelling.
- These models work within the fable framework, which provides the tools to evaluate, visualise, and combine models in a workflow consistent with the tidyverse.

Example

```
p_load(fable, tsibble, tsibbledata, lubridate, dplyr, feast)

aus_retail %>%
  filter(
    State %in% c("New South Wales", "Victoria"),
    Industry == "Department stores"
  ) %>%
  model(
    ets = ETS(box_cox(Turnover, 0.3)),
    arima = ARIMA(log(Turnover)),
    snaive = SNAIVE(Turnover)
  ) %>%
  forecast(h = "2 years") %>%
  autoplot(filter(aus_retail, year(Month) > 2010), level = NULL)
```

Example