

# tsibble and fable



with xaringan

Kevin Song

2020-02-19

1) tsibble

2) fable

# 1) tsibble

```
install.packages("tsibble")
```

■ The **tsibble** package extends the tidyverse to temporal data.

- Built on top of the tibble, **tsibble** is a data- and model-oriented object.
- Comparison to R-based time series objects (**ts**, **zoo** and **xts**)
  - **tsibble** preserves time indices as the essential data column and makes heterogeneous data structures possible.
  - Beyond the tibble-like representation, **key** is introduced to uniquely identify observational units over time (index).
  - The tsibble package aims at managing temporal data and getting analysis done in a fluent workflow.

# Basic Structure

- To coerce data frame to tsibble, we need to declare **key** and **index**.
- **tsibble()** creates a *tsibble object* and **as\_tsibble()** is an S3 method to coerce other objects to a tsibble.
  - Vector/Matrix objects can be automated to a tsibble using **as\_tsibble()** without any specification.
  - For a tibble or data frame, **as\_tsibble()** requires to declare the **key** and **index** variables.

# Basic Structure

■ To coerce data frame to tsibble, we need to declare **key** and **index**.

- In tsibble:
  1. **Key** is a set of variables that define subjects measured over time.
    - **identifier**
  2. **Index** is a variable with inherent ordering from past to present
    - **time-relevant variable**
  3. Each observation should be uniquely identified by **key** and **index**.
  4. Other non-identified variables are referred to as *measured* variables.

# Contextual semantics: **key** and **index**

- **Key** can be comprised of empty, one, or more variables.
  - For multiple variables, declare `key = c(x1, x2, x3)`.
  - can be created in conjunction with tidy selectors like `starts_with()`
- **Index** supports an extensive range of indices
  - native time classes in R (Date, POSIXct, difftime, etc.)
  - tsibble's new additions (yearweek, yearmonth, and yearquarter)
- **Interval** is computed from index based on corresponding time class.

Class	interval
integer/numeric/ordered	either "unit" or "year"(Y)
yearquarter/yearqtr	"quarter"(Q)
yearmonth/yearmon	"month"(M)
yearweek	"week"(W)
Date / difftime	"day"(D) / W, D, h, m, s
POSIXt, hms / nanotime	h, m, s, us, ms / ns

## Ex) nycflights13::weather

- The `weather` data contains the hourly records (temperature, humid and precipitation) over the year of 2013 at three stations (JFK, LGA and EWR).

```
weather <- nycflights13::weather %>%  
  select(origin, time_hour, temp, humid, precip)
```

```
weather
```

```
## # A tibble: 26,115 x 5
```

```
##   origin time_hour      temp humid precip  
##   <chr>   <dtm>      <dbl> <dbl>  <dbl>  
## 1 EWR     2013-01-01 01:00:00  39.0  59.4    0  
## 2 EWR     2013-01-01 02:00:00  39.0  61.6    0  
## 3 EWR     2013-01-01 03:00:00  39.0  64.4    0  
## 4 EWR     2013-01-01 04:00:00  39.9  62.2    0  
## 5 EWR     2013-01-01 05:00:00  39.0  64.4    0  
## 6 EWR     2013-01-01 06:00:00  37.9  67.2    0  
## 7 EWR     2013-01-01 07:00:00  39.0  64.4    0  
## 8 EWR     2013-01-01 08:00:00  39.9  62.2    0  
## 9 EWR     2013-01-01 09:00:00  39.9  62.2    0  
## 10 EWR    2013-01-01 10:00:00  41    59.6    0  
## # ... with 26,105 more rows
```

## Ex) nycflights13::weather

- `time_hour` containing the date-times should be declared as **index**
- `origin` should be declared as **key**.
- Other columns can be considered as measured variables.

```
weather_tsbl <- as_tsibble(weather,  
                             key = origin,  
                             index = time_hour)  
weather_tsbl
```

```
## # A tsibble: 26,115 x 5 [1h] <America/New_York>
```

```
## # Key:          origin [3]
```

	origin	time_hour	temp	humid	precip
	<chr>	<dtm>	<dbl>	<dbl>	<dbl>
## 1	EWR	2013-01-01 01:00:00	39.0	59.4	0
## 2	EWR	2013-01-01 02:00:00	39.0	61.6	0
## 3	EWR	2013-01-01 03:00:00	39.0	64.4	0
## 4	EWR	2013-01-01 04:00:00	39.9	62.2	0
## 5	EWR	2013-01-01 05:00:00	39.0	64.4	0
## 6	EWR	2013-01-01 06:00:00	37.9	67.2	0
## 7	EWR	2013-01-01 07:00:00	39.0	64.4	0
## 8	EWR	2013-01-01 08:00:00	39.9	62.2	0



# fill\_gaps()

| **fill\_gaps()** turns implicit missing values into explicit.

- If the observations are made at regular time interval, we could turn these implicit missingness to be explicit simply using **fill\_gaps()**.
- **fill\_gaps()** also handles filling in **time gaps** by values or functions and respects time zones for date-times.

```
full_weather <- weather_tsbl %>%  
  fill_gaps(precip = 0) %>%  
#filling gaps in precip with 0  
  group_by_key() %>%  
  tidyr::fill(temp, humid, .direction = "down")  
#replaces NAs with its previous observation for each origin (key)  
  
full_weather
```

```
## # A tsibble: 26,190 x 5 [1h] <America/New_York>  
## # Key:      origin [3]  
## # Groups:   origin [3]  
##   origin time_hour      temp humid precip  
##   <chr>   <dtm>         <dbl> <dbl> <dbl>
```

# fill\_gaps()

	origin	time_hour	temp	humid	precip
1	EWR	2013-01-01 01:00:00	39.02	59.37	0
2	EWR	2013-01-01 02:00:00	39.02	61.63	0
3	EWR	2013-01-01 03:00:00	39.02	64.43	0
4	EWR	2013-01-01 04:00:00	39.92	62.21	0
5	EWR	2013-01-01 05:00:00	39.02	64.43	0
6	EWR	2013-01-01 06:00:00	37.94	67.21	0
7	EWR	2013-01-01 07:00:00	39.02	64.43	0
8	EWR	2013-01-01 08:00:00	39.92	62.21	0
9	EWR	2013-01-01 09:00:00	39.92	62.21	0
10	EWR	2013-01-01 10:00:00	41.00	59.65	0
11	EWR	2013-01-01 11:00:00	41.00	57.06	0
12	EWR	2013-01-01 13:00:00	39.20	69.67	0
13	EWR	2013-01-01 14:00:00	39.02	54.68	0
14	EWR	2013-01-01 15:00:00	37.94	57.04	0
15	EWR	2013-01-01 16:00:00	37.04	49.62	0
16	EWR	2013-01-01 17:00:00	35.96	49.83	0



	origin	time_hour	temp	humid	precip
1	EWR	2013-01-01 01:00:00	39.02	59.37	0
2	EWR	2013-01-01 02:00:00	39.02	61.63	0
3	EWR	2013-01-01 03:00:00	39.02	64.43	0
4	EWR	2013-01-01 04:00:00	39.92	62.21	0
5	EWR	2013-01-01 05:00:00	39.02	64.43	0
6	EWR	2013-01-01 06:00:00	37.94	67.21	0
7	EWR	2013-01-01 07:00:00	39.02	64.43	0
8	EWR	2013-01-01 08:00:00	39.92	62.21	0
9	EWR	2013-01-01 09:00:00	39.92	62.21	0
10	EWR	2013-01-01 10:00:00	41.00	59.65	0
11	EWR	2013-01-01 11:00:00	41.00	57.06	0
12	EWR	2013-01-01 12:00:00	41.00	57.06	0
13	EWR	2013-01-01 13:00:00	39.20	69.67	0
14	EWR	2013-01-01 14:00:00	39.02	54.68	0
15	EWR	2013-01-01 15:00:00	37.94	57.04	0
16	EWR	2013-01-01 16:00:00	37.04	49.62	0

# index\_by() + summarise()

| **index\_by()** + **summarise()** to aggregate over calendar periods

- tidy data representation mostly supports part of a “data pipeline” in time-based context.
- **index\_by()** is the counterpart of **group\_by()** in temporal context, but it groups the time index only.
- In conjunction with **index\_by()**, **summarise()** and its scoped variants aggregate variables over calendar periods.
- **index\_by()** goes hand in hand with the index functions including **as.Date()**, **yearweek()**, **yearmonth()** and **yearquarter()**, as well as other friends from **lubridate**.

# index\_by() + summarise()

**Example)** To compute *average* temperature and *total* precipitation *per month*, we can apply `yearmonth()` to the index variable (referred to as `.`).


- While collapsing rows (by `summarise()`), `group_by()` and `index_by()` take care of updating the key and index respectively.

```
full_weather %>%  
  group_by_key() %>%  
  index_by(year_month = ~ yearmonth(.)) %>%  
  # monthly aggregates with new variable name (year_month)  
  summarise(  
    avg_temp = mean(temp, na.rm = TRUE),  
    ttl_precip = sum(precip, na.rm = TRUE)  
  )
```

```
## # A tsibble: 36 x 4 [1M]  
## # Key:      origin [3]  
##   origin year_month avg_temp ttl_precip  
##   <chr>      <mth>      <dbl>      <dbl>  
## 1 EWR      2013 Jan      35.6        3.53  
## 2 EWR      2013 Feb      34.2        3.83  
## 3 EWR      2013 Mar      40.1         3
```

# Rolling with `slide()`, `tile()`, `stretch()`

- Several functions in `tsibble` allow for different variations of moving windows using `purrr`-like syntax:
  - `slide()`: sliding window with overlapping observations.
  - `tile()`: tiling window without overlapping observations.
  - `stretch()`: fixing an initial window and expanding to include more observations.



## 2) fable

```
install.packages("fable")
```

**fable** package provides a collection of commonly used univariate and multivariate time series forecasting models

- Forecasting models include exponential smoothing via state space models and automatic ARIMA modelling.
- These models work within the fable framework, which provides the tools to evaluate, visualise, and combine models in a workflow consistent with the tidyverse.

# Example

```
p_load(fable, tsibble, tsibbledata, lubridate, dplyr, feast)

aus_retail %>%
  filter(
    State %in% c("New South Wales", "Victoria"),
    Industry == "Department stores"
  ) %>%
  model(
    ets = ETS(box_cox(Turnover, 0.3)),
    arima = ARIMA(log(Turnover)),
    snaive = SNAIVE(Turnover)
  ) %>%
  forecast(h = "2 years") %>%
  autoplot(filter(aus_retail, year(Month) > 2010), level = NULL)
```

Sorry, I DO NOT know what `ets` and `snaive` are.

- Visit [Forecasting: Principles and Practice](#)
  - looks super helpful when dealing with timeseries data

# Example